

AWS Certified AI Practitioner AIF-C01

Section 7: Amazon Q

Amazon Q is AWS's family of **generative AI assistants** designed to help users perform tasks using natural language. Instead of searching through documentation, writing code manually, or analyzing large datasets yourself, Amazon Q allows you to ask questions in plain English and receive intelligent responses based on AWS services, enterprise data, or software development workflows. Amazon Q is built using **Amazon Bedrock** and multiple foundation models while incorporating enterprise-grade security, permissions, and governance. It is not a single product but rather a collection of AI-powered assistants designed for different use cases.

Amazon Q Family

The Amazon Q ecosystem consists primarily of four products:

1. Amazon Q
2. Amazon Q Business
3. Amazon Q Developer
4. Amazon Q Apps

Think of them as different specialists within the same AI family.

Product	Primary Users	Main Purpose
Amazon Q	General AWS users	AI assistant for AWS and cloud tasks
Amazon Q Business	Business employees	Enterprise knowledge assistant
Amazon Q Developer	Software developers	Coding and software development assistant
Amazon Q Apps	Business users	Create AI-powered applications without coding

1. Amazon Q

Amazon Q is a generative AI assistant from AWS that helps users perform cloud-related tasks using natural language. It assists with AWS services, architecture recommendations, troubleshooting, operational guidance, documentation, and automation.

It understands AWS documentation, best practices, and cloud infrastructure.

Simple Explanation

Imagine Amazon Q as an AWS expert sitting next to you.

Instead of asking:

"How do I create an EC2 Auto Scaling Group?"

You simply type:

"Create a highly available web application architecture."

Amazon Q explains everything, generates commands, and even provides architecture recommendations.

What Amazon Q Can Do

- Explain AWS services
- Recommend architectures
- Generate CLI commands
- Generate CloudFormation templates
- Help troubleshoot AWS errors
- Answer AWS documentation questions
- Explain IAM policies
- Recommend security best practices
- Optimize cloud costs
- Explain AWS Well-Architected Framework

Example

You ask:

"Why can't my EC2 instance access S3?"

Amazon Q may respond:

- Check IAM role permissions
- Verify bucket policy
- Confirm VPC endpoint configuration
- Ensure network connectivity
- Check region mismatch

Instead of searching documentation manually, Amazon Q analyzes the situation and provides guidance.

2. Amazon Q Business

Amazon Q Business is an enterprise generative AI assistant that securely connects to an organization's internal data sources, enabling employees to ask questions using natural language and receive answers based on company knowledge while respecting existing permissions.

It functions as an intelligent enterprise search and knowledge assistant.

Simple Explanation

Imagine a company has information stored in:

- SharePoint
- Confluence
- Salesforce
- Google Drive
- Microsoft 365
- Jira
- Amazon S3
- Internal databases

Instead of opening each application individually, employees ask Amazon Q Business:

"What is our vacation policy?"

Amazon Q Business searches all connected systems and provides the answer.

Key Features:

Enterprise Search

Search across multiple enterprise applications simultaneously.

Permission-Aware Responses

Amazon Q Business only displays information the user is authorized to access.

For example:

- HR documents are visible only to HR staff.
- Finance reports are accessible only to Finance users.

Natural Language Questions

Instead of using keywords, users ask conversational questions.

Example:

"What are the latest sales numbers for Europe?"

Data Connectors

Amazon Q Business integrates with services such as:

- Microsoft SharePoint
- Microsoft Teams
- Jira
- Confluence
- Salesforce
- Slack
- Google Drive
- Amazon S3

Summarization

It can summarize:

- Reports
- PDFs
- Emails
- Meeting notes
- Policies

Real-World Example

A new employee asks:

"How do I request parental leave?"

Amazon Q Business searches HR documents and responds with the relevant company policy.

3. Amazon Q Developer

Amazon Q Developer is an AI-powered software development assistant that helps developers write, understand, test, document, refactor, and optimize code while assisting with AWS application development and troubleshooting.

Simple Explanation

Amazon Q Developer acts like an experienced senior software engineer available at any time.

It helps with:

- Writing code
- Explaining code
- Debugging
- Reviewing code
- Creating tests
- Optimizing performance

Main Capabilities:

Code Generation

Example prompt:

"Write a Python Lambda function that uploads files to S3."

Amazon Q Developer generates the code.

Code Explanation

Paste complex code and ask:

"Explain this function."

It provides a step-by-step explanation.

Code Refactoring

It improves code readability and performance.

Example:

Convert nested loops into cleaner Python code.

Unit Test Generation

Generate unit tests automatically.

Documentation Generation

Automatically create:

- Function comments
- README files
- API documentation

AWS Assistance

Amazon Q Developer can help create:

- Lambda functions
- IAM policies
- CloudFormation templates
- CDK code
- Terraform snippets
- AWS CLI commands

Security Scanning

Detects:

- SQL injection risks
- Hardcoded secrets
- Insecure code patterns

Supported IDEs

Amazon Q Developer integrates with:

- Visual Studio Code
- JetBrains IDEs
- AWS Cloud9

It also integrates with the AWS Management Console.

Example

Prompt:

"Create a REST API using Lambda, API Gateway, and DynamoDB."

Amazon Q Developer generates the required code and infrastructure guidance.

4. Amazon Q Apps

Amazon Q Apps enables business users to create lightweight AI-powered applications using natural language without traditional software development.

These applications automate workflows, summarize documents, analyze data, and perform repetitive business tasks.

Simple Explanation

Suppose an HR manager wants an application that:

- Reviews resumes
- Summarizes candidates
- Suggests interview questions

Instead of hiring developers, they ask Amazon Q:

"Create an app that reviews resumes."

Amazon Q Apps generates the application.

Key Features:

No-Code Development

No programming knowledge is required.

Natural Language App Creation

Example:

"Create an expense approval application."

Amazon Q Apps builds it.

AI Workflow Automation

Automate tasks such as:

- Report generation
- Email drafting
- Document summarization
- Data extraction
- Customer support responses

Sharing

Applications can be shared with coworkers.

Enterprise Data Access


Apps can use organizational knowledge through Amazon Q Business while respecting permissions.

Example

A finance department creates an app that:

- Reads invoices
- Extracts totals
- Identifies vendors
- Generates summaries

No coding is required.



Amazon Q Business integrates with AWS IAM Identity Center

Amazon Q Business integrates with AWS IAM Identity Center to securely manage employee access to the Q Business application. IAM Identity Center connects users and groups from an organization's identity source—such as Microsoft Entra ID, Okta, Active Directory, or its built-in directory—and provides centralized sign-in and access assignment. When a user asks Amazon Q Business a question, their identity and permissions are used to ensure they can access only the enterprise content they are authorized to view. Administrators can also assign Amazon Q Business subscriptions to specific users or groups through this integration.

Amazon QuickSight

Amazon QuickSight is AWS's fully managed **business intelligence (BI) and data visualization service** that enables organizations to analyze data, create interactive dashboards, and generate reports. It connects to various data sources such as Amazon S3, Amazon Redshift, Amazon RDS, Athena, and third-party databases, allowing users to visualize trends, monitor key performance indicators (KPIs), and gain actionable insights. QuickSight also incorporates generative AI capabilities through **Amazon Q in QuickSight**, enabling users to ask questions

in natural language, automatically generate charts, summaries, and insights without requiring advanced SQL or data analysis expertise. It is designed to scale to thousands of users while offering a pay-per-session pricing model for readers.

Amazon Q for Amazon EC2

Amazon Q for Amazon EC2 is an AI-powered assistant integrated into the Amazon EC2 console that helps users manage, troubleshoot, and optimize EC2 instances using natural language. Instead of manually investigating issues, users can ask questions such as *"Why is my EC2 instance unreachable?"* or *"Why is CPU utilization high?"* Amazon Q analyzes relevant instance information, including configuration, metrics, logs, and AWS best practices, to identify potential causes and recommend solutions. It can also explain EC2 features, suggest performance and cost optimizations, and provide step-by-step guidance for resolving operational issues, helping administrators and developers manage EC2 environments more efficiently.

AWS Chatbot

AWS Chatbot is an AWS service that enables you to receive **AWS notifications and interact with AWS resources directly from collaboration platforms such as Slack and Microsoft Teams**. It integrates with services like **Amazon SNS, Amazon CloudWatch, AWS Health, and AWS CloudFormation** to send alerts about alarms, operational events, and infrastructure changes to chat channels. With appropriate IAM permissions, users can also run approved AWS CLI commands from within the chat application to investigate issues or perform operational tasks without logging into the AWS Management Console. AWS Chatbot helps improve collaboration, incident response, and operational efficiency by bringing AWS monitoring and management capabilities into team communication tools.

Exam Tip (AWS AI Practitioner): Despite its name, **AWS Chatbot is *not* a generative AI service**. It is a **DevOps and notification service** for chat platforms. If a question asks for an AI-powered conversational assistant, the correct answer is typically **Amazon Q**, not AWS Chatbot.

AWS Glue

AWS Glue is a **fully managed, serverless data integration service** that helps you **discover, prepare, transform, clean, and move data** for analytics, machine learning, and application development. It eliminates the need to provision or manage servers, allowing you to focus on building data pipelines rather than infrastructure.

In simple terms, AWS Glue is the **ETL (Extract, Transform, Load) service** of AWS. It extracts data from multiple sources, transforms it into the required format (such as cleaning, filtering, or joining datasets), and loads it into a destination like Amazon S3, Amazon Redshift, or another data store.

Why Do We Need AWS Glue?

Organizations often store data in different formats and locations:

- Customer records in Amazon RDS
- Sales data in Amazon S3
- Logs in Amazon CloudWatch
- Data warehouse in Amazon Redshift

Before this data can be analyzed, it usually needs to be:

- Combined
- Cleaned

- Converted into a common format
- Loaded into a central repository

AWS Glue automates these tasks.

ETL Process

AWS Glue primarily performs **ETL (Extract, Transform, Load)** operations.

1. Extract

Collect data from one or more sources.

Examples:

- Amazon S3
- Amazon RDS
- DynamoDB
- Redshift
- JDBC databases

2. Transform

Modify the data by:

- Removing duplicates
- Filtering unwanted records
- Joining datasets
- Converting formats (CSV → Parquet)
- Renaming columns
- Cleaning missing values

3. Load

Store the transformed data into a destination such as:

- Amazon S3
- Amazon Redshift
- Amazon Aurora
- Another database

Main Components of AWS Glue

1. Glue Crawlers

A **Crawler** automatically scans your data source and discovers:

- Tables
- Columns
- Data types
- Partitions
- File formats

It stores this metadata in the **Glue Data Catalog**.

Example

Suppose you upload this CSV file to Amazon S3:

Name,Age,Salary

John,30,60000

Alice,25,50000

The crawler identifies:

- Table: Employees
- Name → String
- Age → Integer
- Salary → Integer

No manual schema definition is required.

2. Glue Data Catalog

The **Glue Data Catalog** is a centralized metadata repository.

It stores information about datasets, such as:

- Table names
- Column names
- Data types
- Partition details
- File locations

Services like **Amazon Athena**, **Amazon EMR**, **Amazon Redshift Spectrum**, and **Lake Formation** use the Data Catalog to understand the structure of data.

Think of it as:

A library catalog that tells you **what data exists, where it is located, and how it is structured**—it does not store the actual data.

3. Glue ETL Jobs

Glue Jobs perform the actual data transformation.

They can:

- Filter records
- Join tables
- Convert formats
- Clean data
- Aggregate data
- Enrich datasets

Glue jobs are typically written in:

- Python (PySpark)
- Scala

AWS Glue can also generate ETL scripts automatically based on your configuration.

4. Glue Studio

Glue Studio provides a **visual drag-and-drop interface** for creating ETL pipelines.


Benefits:

- Little or no coding
- Easy workflow design
- Automatic code generation
- Pipeline monitoring

5. Glue Workflows

Workflows allow you to orchestrate multiple Glue jobs and crawlers.

Example:



6. Glue Triggers

Triggers determine when a job starts.

They can be:

- On a schedule
- On-demand
- Event-driven
- After another job completes

Common Data Sources

AWS Glue supports:

- Amazon S3
- Amazon RDS
- Amazon Aurora
- Amazon Redshift
- DynamoDB
- JDBC-compatible databases
- Data lakes

Common Data Destinations

Glue can write data to:

- Amazon S3

- Amazon Redshift
- Amazon Aurora
- JDBC databases
- Data lakes

Serverless Nature

AWS Glue is **serverless**, meaning:

- No EC2 instances to manage
- No cluster provisioning
- Automatic scaling
- Pay only for the resources used during job execution

Integration with Other AWS Services

Service	How Glue Integrates
Amazon S3	Read and write datasets
Amazon Athena	Query data using the Glue Data Catalog
Amazon Redshift	Load transformed data into a data warehouse
AWS Lake Formation	Use the Glue Data Catalog and enforce data access controls
Amazon EMR	Share metadata from the Glue Data Catalog
Amazon SageMaker AI	Prepare and transform training data for ML models
Amazon QuickSight	Visualize cleaned and transformed data

PartyRock

PartyRock is a **free, browser-based generative AI application playground from AWS** that allows users to quickly build, experiment with, and share AI-powered applications **without writing code**. It is designed for learning and rapid prototyping rather than production deployments. PartyRock uses **Amazon Bedrock** and its underlying **foundation models** to generate text, summarize content, answer questions, create images (depending on the model), and perform other generative AI tasks through simple prompts and drag-and-drop widgets.

In simple terms, PartyRock is a **no-code environment** where anyone can build small generative AI applications by describing what they want in natural language. For example, you can create an app that generates interview questions, summarizes articles, creates travel itineraries, or drafts marketing content in just a few minutes. Users can also remix existing PartyRock apps created by others, making it an excellent platform for learning prompt engineering and exploring generative AI concepts.

Key Features

- **No coding required** – Build AI apps using prompts and visual widgets.
- **Powered by Amazon Bedrock** – Uses foundation models available through Amazon Bedrock.
- **Rapid prototyping** – Quickly test and experiment with AI ideas.
- **Share and remix apps** – Publish apps or customize apps created by the community.
- **Educational platform** – Ideal for learning prompt engineering and generative AI concepts.

For the **AWS Certified AI Practitioner (AIF-C01)** exam:

- If the question mentions **learning, experimenting, or creating simple AI apps without coding**, the answer is **PartyRock**.
- If the question mentions **building scalable, production-ready generative AI applications using APIs and foundation models**, the answer is **Amazon Bedrock**.

AWS AI Practitioner Practice Questions

1. A company wants employees to ask natural-language questions about internal HR policies, technical documentation, and company procedures. The solution must respect each employee's access permissions. Which service is MOST appropriate?

- A. Amazon Bedrock
- B. Amazon Q Business
- C. Amazon SageMaker AI
- D. AWS Glue

Answer: B — Amazon Q Business

Explanation: Amazon Q Business is a generative AI assistant for enterprise knowledge. It can connect to organizational data sources and provide permission-aware answers based on the user's authorized content.

2. What is the primary purpose of Amazon Q Developer?

- A. Train custom machine learning models
- B. Build business intelligence dashboards
- C. Assist developers with software development and AWS tasks
- D. Perform enterprise ETL operations

Answer: C — Assist developers with software development and AWS tasks

Explanation: Amazon Q Developer provides AI-powered assistance for coding, code explanation, debugging, testing, AWS development, and related developer workflows.

3. A developer wants to build a custom generative AI chatbot using foundation models from multiple model providers without managing the underlying model infrastructure. Which AWS service should the developer use?

- A. Amazon Q Business
- B. Amazon Bedrock
- C. Amazon Kendra
- D. Amazon QuickSight

Answer: B — Amazon Bedrock

Explanation: Amazon Bedrock provides managed access to foundation models and capabilities for building generative AI applications without managing the underlying infrastructure.

4. What role do foundation models play in generative AI?

- A. They are pre-trained models that can be adapted or used for many downstream AI tasks
- B. They exclusively store enterprise documents
- C. They manage IAM users and groups
- D. They only perform ETL operations

Answer: A — They are pre-trained models that can be adapted or used for many downstream AI tasks

Explanation: Foundation models are trained on large amounts of data and can support tasks such as generation, summarization, question answering, reasoning, and multimodal processing depending on the model.

5. A data scientist needs to build, train, tune, and deploy a custom machine learning model for predicting customer churn. Which AWS service is MOST appropriate?

- A. Amazon Q
- B. Amazon SageMaker AI
- C. Amazon Kendra
- D. PartyRock

Answer: B — Amazon SageMaker AI

Explanation: SageMaker AI provides tools for the ML lifecycle, including data preparation, model development, training, tuning, deployment, and monitoring.

6. Which statement BEST describes the difference between Amazon Q and Amazon Bedrock?

- A. Amazon Q is primarily a ready-to-use AI assistant, while Amazon Bedrock provides a platform for building generative AI applications
- B. Amazon Q trains ML models, while Bedrock performs ETL
- C. Amazon Q is a database, while Bedrock is a data warehouse
- D. Amazon Q and Bedrock provide exactly the same functionality

Answer: A

Explanation: Think Q = use an AI assistant and Bedrock = build generative AI applications.

7. A company wants an intelligent search solution that allows employees to find information across enterprise documents using natural-language queries. Which AWS service is specifically designed for intelligent enterprise search?

- A. Amazon Kendra
- B. AWS Glue
- C. Amazon SageMaker AI
- D. Amazon EC2

Answer: A — Amazon Kendra

Explanation: Amazon Kendra is an intelligent enterprise search service designed to search and retrieve information from organizational content.

8. What is the primary role of AWS IAM Identity Center when integrated with Amazon Q Business?

- A. Train foundation models
- B. Transform enterprise data
- C. Manage workforce identities and access to Q Business applications
- D. Create QuickSight dashboards

Answer: C — Manage workforce identities and access

Explanation: IAM Identity Center provides centralized workforce access. In Q Business environments, it helps organizations control which users and groups can access applications and associated resources.

9. A business user wants to create a lightweight generative AI application using company knowledge without writing traditional application code. Which capability is designed for this use case?

- A. Amazon Q Apps
- B. Amazon EC2
- C. Amazon Kendra
- D. AWS Glue Crawler

Answer: A — Amazon Q Apps

Explanation: Amazon Q Apps enables users to create lightweight generative AI-powered applications from natural-language descriptions using organizational knowledge available through Amazon Q Business.

10. What is the primary purpose of Amazon QuickSight?

- A. Train foundation models
- B. Business intelligence, analytics, and data visualization
- C. Software development assistance
- D. Manage IAM policies

Answer: B — Business intelligence, analytics, and data visualization

Explanation: QuickSight is AWS's managed BI service for analyzing data and creating interactive visualizations, reports, and dashboards.

11. A company has raw CSV files in Amazon S3. It wants to automatically discover the files' schema and make the metadata available for analytics. Which AWS Glue component should it use?

- A. Glue Studio
- B. Glue Crawler
- C. Glue Trigger
- D. Glue ETL Job

Answer: B — Glue Crawler

Explanation: A Glue Crawler scans data sources, identifies schemas and partitions, and creates or updates metadata in the AWS Glue Data Catalog.

12. What does the AWS Glue Data Catalog primarily store?

- A. The actual business data
- B. Foundation models
- C. Metadata describing datasets
- D. EC2 snapshots

Answer: C — Metadata describing datasets

Explanation: The Glue Data Catalog stores metadata such as table definitions, schemas, columns, partitions, and data locations. The actual data can remain in S3 or other data stores.

13. A company needs to extract sales data, remove duplicate records, convert CSV files to Parquet, and store the processed data in Amazon S3. Which AWS service is MOST appropriate?

- A. AWS Glue
- B. Amazon Q Developer
- C. Amazon Kendra
- D. IAM Identity Center

Answer: A — AWS Glue

Explanation: This is a classic **ETL/data integration** workload. Glue can extract, clean and transform the data, then write the transformed output to S3.

14. What does ETL stand for?

- A. Evaluate, Train, Learn
- B. Extract, Transform, Load
- C. Execute, Transfer, Log
- D. Extract, Train, Launch

Answer: B — Extract, Transform, Load

Explanation: ETL means extracting data from sources, transforming it into the required form, and loading it into a destination.

15. A beginner wants to experiment with generative AI by creating a simple application without writing code or provisioning AWS infrastructure. Which AWS offering is designed for this learning and experimentation use case?

- A. Amazon SageMaker AI
- B. PartyRock
- C. Amazon Redshift
- D. AWS Glue

Answer: B — PartyRock

Explanation: PartyRock is a generative AI playground from AWS designed for experimenting with and building simple generative AI applications without traditional coding.

16. Which statement BEST describes PartyRock compared with Amazon Bedrock?

- A. PartyRock is oriented toward no-code experimentation, while Bedrock is a managed service for building generative AI applications
- B. PartyRock is an enterprise database, while Bedrock is a data warehouse
- C. PartyRock performs ETL, while Bedrock performs SQL queries
- D. PartyRock manages IAM identities, while Bedrock manages AWS accounts

Answer: A

Explanation: For exam purposes, associate **PartyRock with learning/prototyping** and **Amazon Bedrock with developing generative AI applications using foundation models and APIs**.

17. Which Amazon Q offering is MOST appropriate for a developer who wants AI assistance with generating, explaining, debugging, and improving code?

- A. Amazon Q Business
- B. Amazon Q Developer
- C. Amazon Q Apps
- D. Amazon QuickSight

Answer: B — Amazon Q Developer

Explanation: Amazon Q Developer is specifically designed for software-development and AWS-development workflows.

18. Which service combination BEST fits a pipeline where raw data in S3 is cataloged, transformed, queried, and then visualized?

- A. AWS Glue → Amazon Athena → Amazon QuickSight
- B. Amazon Q Developer → EC2 → IAM Identity Center
- C. Amazon Kendra → Lambda → PartyRock
- D. Amazon Q Apps → IAM → CloudFormation

Answer: A — AWS Glue → Amazon Athena → Amazon QuickSight

Explanation: Glue can catalog and transform the data, Athena can query data in S3 using SQL, and QuickSight can visualize the resulting data through dashboards and reports.

19. Which scenario BEST fits Amazon Q Business rather than Amazon Q Developer?

- A. Generating unit tests for a Python application
- B. Explaining source code in an IDE
- C. Answering employee questions using internal company policies and documents
- D. Helping a developer understand an AWS SDK

Answer: C — Answering employee questions using internal company information

Explanation: Q Business = enterprise knowledge and productivity. Q Developer = software development and AWS development assistance.

20. A company wants employees to receive AWS operational notifications in team collaboration channels. Which AWS capability historically addressed this use case?

- A. AWS Chatbot
- B. Amazon SageMaker AI
- C. AWS Glue
- D. Amazon Kendra

Answer: A — AWS Chatbot

RAPID REVISION SHEET

AMAZON Q FAMILY

Amazon Q → Generative AI assistant family

Q Business → Enterprise knowledge + employee assistant

Q Developer → Coding + AWS development assistant
Q Apps → Build lightweight GenAI apps using natural language
IAM Identity Center → Workforce identity/access for Q Business

GENERATIVE AI / ML
Amazon Bedrock → Build GenAI applications using foundation models
SageMaker AI → Build/train/deploy/manage ML models
PartyRock → No-code GenAI learning and experimentation

SEARCH
Amazon Kendra → Intelligent enterprise search

DATA / ANALYTICS
AWS Glue → Serverless data integration / ETL
Glue Crawler → Discover schemas
Glue Data Catalog → Store metadata about datasets
Amazon Athena → Query S3 data using SQL
Amazon QuickSight → BI, dashboards, analytics, visualization

MEMORY TRICKS
Q Business = ASK your BUSINESS data
Q Developer = ASSIST the DEVELOPER
Q Apps = CREATE lightweight AI apps
Bedrock = BUILD GenAI applications
SageMaker = BUILD / TRAIN / DEPLOY ML models
Kendra = SEARCH enterprise information
Glue = PREPARE / TRANSFORM data
QuickSight = VISUALIZE data
PartyRock = EXPERIMENT with GenAI

Amazon Q vs Amazon Bedrock vs Amazon SageMaker AI vs Amazon Kendra

These four services are frequently compared because they all relate to AI, but they solve different problems.

1. Amazon Bedrock: Build generative AI applications using managed foundation models.

Best For

- Chatbots
- AI assistants
- Text generation
- Image generation
- Summarization

- RAG (Retrieval-Augmented Generation)
- AI APIs

You Use Bedrock When

You want to build your own ChatGPT-like application without managing model infrastructure.

2. Amazon Q:

Ready-to-use AI assistant for AWS users, developers, and businesses.

You Use Amazon Q When

You want an AI assistant rather than building one yourself.

3. Amazon SageMaker AI

Amazon SageMaker AI is a managed machine learning platform for preparing data, building, training, tuning, deploying, and monitoring machine learning models at scale.

Unlike Amazon Bedrock, which focuses on foundation models, SageMaker AI supports the complete machine learning lifecycle and allows you to work with traditional ML as well as generative AI workflows.

4. Amazon Kendra

Amazon Kendra is an intelligent enterprise search service that helps users find information across multiple content repositories using natural language search.

It specializes in **search and retrieval**, not generative text creation.

Example

Ask:

"Where is the HR handbook?"

Kendra finds the document.

It does **not** write a new answer like Amazon Q Business can.

Decision Flow

